

Final Exam

- Do not open this quiz booklet until directed to do so. Read all the instructions on this page.
- When the quiz begins, write your name on every page of this quiz booklet.
- You have **180 minutes to earn 180 points**. Do not spend too much time on any one problem. Read them all first, and attack them in the order that allows you to make the most progress.
- **You are allowed a 3-page cheat sheet.** No calculators or programmable devices are permitted. No cell phones or other communications devices are permitted.
- Write your solutions in the space provided. If you need more space, use **the back of the sheet** containing the problem. Pages may be separated for grading.
- Do not waste time and paper rederiving facts that we have studied. Simply cite them.
- When writing an algorithm, a **clear** description in English will suffice. Pseudo-code is not required unless you find that it helps with the clarity of your presentation.
- **Pay close attention to the instructions for each problem.** Depending on the problem, partial credit may be awarded for incomplete answers.

Problem	Parts	Points	Grade	Grader
0	2	2		
1	9	18		
2	3	20		
3	2	15		
4	3	15		
5	1	20		
6	1	20		
7	1	20		
8	2	20		
9	2	30		
Total		180		

Name: _____

Circle your recita- tion:	R01 Daniel Manesh 10AM	R02 Jennifer Jang 10AM	R03,6 Adam Yedidia 11,12PM	R04,5 Kevin Chen 11,12PM	R07 Casey O'Brien 1PM	R08 Jonathan Chien 2PM	R09 Victor Valov 3PM
------------------------------------	----------------------------------------	----------------------------------------	--------------------------------------------	------------------------------------------	---------------------------------------	----------------------------------------	--------------------------------------

Problem 0. What is Your Name? [2 points] (2 parts)

(a) [1 point] Flip back to the cover page. Write your name there.

(b) [1 point] Flip back to the cover page. Circle your recitation section.

Problem 1. True or False [18 points]

For each of the following questions, circle either T (True) or F (False). There is no need to justify the answers.

Each correct answer is worth 2 points and each incorrect answer receives -2 points. If the total is negative, you will receive 0 points for this problem.

- (a) **T F** Suppose algorithm FAST runs in $\Theta(n^2)$ and algorithm SLOW runs in $\Theta(2^n)$. FAST will run faster than SLOW on any input.

Solution: False. FAST can have high constant overhead.

- (b) **T F** It is possible to merge k sorted lists of objects, each of size n , into a single sorted list of size nk in $O(nk)$ time using subroutines that compare a given pair and copy a given object in one time step.

Solution: False. It takes $\Omega(nk \log k)$. Otherwise we could perform comparison-based sorting in $o(n \log n)$ time.

- (c) **T F** A balanced binary search tree can be augmented to implement all priority queue operations (MIN, EXTRACT-MIN, INCREASE-KEY, INSERT) at least as fast (asymptotically) as a min-heap.

Solution: True. Augment with a pointer to the minimum.

- (d) **T F** For all positive integers a, b , and m , we have $(a \bmod m)(b \bmod m) = ab \bmod m$.

- (e) **T F** For all odd primes p such that $\gcd(3, p-1) = 1$, and for all $a \in \mathbb{Z}_p^*$, if x is the inverse of 3 modulo $p-1$, then a^x is a cube root of a modulo p .

Solution: True. Write $x = (p-1)t + 3$ for some integer t and use the fact that $a^{p-1} = 1 \pmod{p}$.

- (f) **T F** For all primes p that end in 99 (such as 199), if a is a square modulo p , then $\sqrt{a} = a^{(p+1)/2} \pmod{p}$.

Solution: False. However, since $p \equiv 3 \pmod{4}$, we have $\sqrt{a} = \pm a^{(p+1)/4} \pmod{p}$.

- (g) **T F** For a directed graph, the absence of back edges with respect to a DFS tree implies that the graph is acyclic.

- (h) **T F** At least one of the following statements is in NP: 1) Given undirected graphs G and H , the graphs are isomorphic; 2) Given undirected graphs G and H , the graphs are not isomorphic.

- (i) **T F** There is a polynomial time verifiable NP proof for the following statement: "Given positive integers n and k such that $1 < k < n$, there is no integer m such that $1 < m \leq k$ and n is divisible by m ."

Problem 2. Short answers [20 points]

- (a) [5 points] You are given a weighted, directed graph $G = (V, E)$ with n vertices and positive integer edge weights. Given a node t and an integer W , design an algorithm running in time $O(n^2 \log n)$ which determines, for every $s \in V$, if there exists a path of weight less than W from s to t .

Solution: Reverse the direction of the edges on the graph, then run Dijkstra's Algorithm. This will get you the shortest path from each node s to our destination t . If any of these path lengths are greater than W , return False.

- (b) [5 points] Give an algorithm with strictly faster asymptotic running time than the trivial $\Theta(n)$ for the following problem: Given an array of n integers $A = [a_1, a_2, \dots, a_n]$ such that the number of inversions, defined as pairs of indices (i, j) where $i < j$ but $a_i > a_j$, is at most $\log_2 n$, and given an integer x , decide whether x is an element of A .

Solution: Let's define the number of inversions as k . Do a binary search, but instead of just checking the middle element, check a window of size $2k$ around the middle element. For each iteration, if x is not found in this window, we can safely discard half of the list. This algorithm is $O(k \lg(n)) = O(\lg^2(n))$

- (c) [10 points] Given an undirected weighted graph $G = (V, E)$ with non-negative weights, and two vertices $s, t \in V$, design an efficient algorithm that returns the shortest-weight path from s to t among all possible paths that take no more than 5 edges; or returns `None` if no such path exists. Full credit will be given to algorithms that run in time $O(n^2 \log n)$, where n is the number of vertices.

Solution: We create a graph $G' = (V', E')$ as follows. For each $v \in V$, create nodes v_1, v_2, v_3, v_4, v_5 in V' . For each $(u, v) \in E$, create edge (u_i, v_{i+1}) for $1 \leq i \leq 4$. Then, run Dijkstra on G' . The shortest path is the one corresponding to $\min\{\delta(v_{B_1}, v_{L_1}), \delta(v_{B_1}, v_{L_2}), \delta(v_{B_1}, v_{L_3}), \delta(v_{B_1}, v_{L_4}), \delta(v_{B_1}, v_{L_5})\}$.

Problem 3. Dynamic programming [15 points]

You are given a **rooted tree** $G = (V, E)$, where for each vertex v , there is an associated positive integer reward r_v . The goal is to compute $S \subseteq V$ which maximizes the total reward, but which also does not contain any adjacent nodes. That is, we want to choose $S = \{v_1, v_2, \dots, v_k\}$ to maximize $\sum_{v \in S} r_v$ with the constraint that if (u, v) is an edge in our graph, then S can't contain both u and v .

You decide to solve this problem using Dynamic Programming, using the following subproblems:

$R[v, True] :=$ The maximum total reward from the subtree rooted at v , where v is included in S .

$R[v, False] :=$ The maximum total reward from the subtree rooted at v , where v is **not** included in S .

- (a) [5 points] Write down the values for $R[v, True]$ and $R[v, False]$ when v is a leaf.

$$R[v, True] =$$

$$R[v, False] =$$

- (b) [10 points] Write down the recurrence for $R[v, True]$ and $R[v, False]$ when v is **not** a leaf.

$$R[v, True] =$$

$$R[v, False] =$$

Solution: We have:

$$\begin{aligned} R(v, True) &= \sum_{w \in \text{children}(v)} R(w, False) + r(v) \\ R(v, False) &= \sum_{w \in \text{children}(v)} \max(R(w, True), R(w, False)) \end{aligned}$$

Problem 4. Hashing [15 points]

Ben Bitdiddle is investigating the performance of different hash table collision strategies. His universe of possible inputs is $U = \{0, 4, 8, 12, 16, \dots, 100\}$. His hash function is

$$h(k) = k \bmod m,$$

where m is the size of the table.

- (a) [5 points] Ben inserts the specific elements 32, 12, 4, and 76 into the table (in the specified order). Indicate the contents of each entry in a hash table of size 7 after inserting the above elements with chaining as the collision resolution strategy.

Hash Entry	Key(s)
0	
1	
2	
3	
4	
5	
6	

- (b) [5 points] Repeat part (a), but this time use linear probing as the collision resolution strategy.

Hash Entry	Key(s)
0	
1	
2	
3	
4	
5	
6	

Solution: Chaining: 4: [32, 4], 5: [12], 6: [76], rest empty Linear probing: 0: [76], 4: [32], 5: [12], 6: [4] Double hashing: 2: [76], 4: [32], 5: [12], 6: [4]

- (c) [5 points] Ben notices the load factor of the table is large and proposes resizing the table to size 20. Is this a good choice of table size considering the elements that may be inserted? Explain why or why not.

Solution: No. m and every possible input share a factor of 4, meaning the effective new size of the hash table will only be 5.

Problem 5. Hanging Posters [20 points]

A company is hanging advertisement posters in a long hallway. There are n possible distinct locations x_1, \dots, x_n for hanging posters where each x_i is an integer that specifies the distance of the i th location (in meters) from the entrance of the hallway. For each $i \in \{1, \dots, n\}$, hanging a poster at location i brings a revenue of r_i for the company, where r_i is a positive integer. In order to maximize the revenue, the company desires to hang as many posters as possible, however, there is the limiting constraint that no two posters can be within five meters of each other.

Design an efficient algorithm that takes the locations x_1, \dots, x_n (in no particular order) and returns the maximum total revenue that can be obtained from choosing any subset of the locations. Full credit will be given to algorithms that run in $O(n \log n)$.

For example, suppose that $n = 4$, and $(x_1, x_2, x_3, x_4) = (6, 8, 14, 12)$, and $(r_1, r_2, r_3, r_4) = (5, 6, 1, 5)$. Then, the optimal solution is to place posters at x_1 and x_4 for a total revenue of $5 + 5 = 10$.

Problem 6. Palindromes [20 points]

Given a string $x = x_1, \dots, x_n$, design an efficient algorithm to find the minimum number of characters that need to be inserted to make it a palindrome (recall that a palindrome is a string such as “racecar” that reads the same backwards). Analyze the running time of your algorithm and justify its correctness.

For example, when $x = \text{“ab3bd”}$, we need to insert two characters (one “d” and one “a”), to get either the of the palindromes “dab3bad” or “adb3bda”.

Problem 7. Housekeeping [20 points]

Ben Bitdiddle likes to make a little extra money doing tasks around the house. Each task has an associated time t that it will take Ben to do the task, and a value v , representing the amount of money that Ben will be paid to do it.

Ben wants to design a data structure to maintain the set of tasks he can perform. Whenever someone thinks of a new task for Ben to do, Ben inserts the task into the structure. Whenever Ben has some amount of free time t , he wants to query the structure to determine the maximum value he can earn from a single task in time at most t . Assume that Ben never needs to remove a task from the structure.

Specifically, design a data structure which supports the following operations:

- **INSERT**(t, v): Insert a task into the structure which takes time t and has value v .
- **QUERY**(t): Return the maximum value that Ben can get from performing a single task in time at most t .

Both operations should run in $O(\log n)$, where n is the number of tasks in the structure at the time that the operation is called. For this problem, **no** partial credit will be awarded for slower solutions.

Solution: Maintain a BST keyed by time. Augment each node with the maximum value in the subtree rooted at that node. Insertion is just normal BST insertion, updating the augmentation as the node is inserted. To perform a query, just search for t in the BST. As we go down, we take the maximum of all the values of nodes less than t , which we can easily do with our augmentation.

Problem 8. Congruence Systems [20 points]

- (a) [10 points] Prove or disprove the following: There is a positive integer x satisfying both of the following equations.

$$\begin{aligned}x &\equiv 1 \pmod{6} \\x &\equiv 2 \pmod{15}.\end{aligned}$$

Solution: False. Because assuming such an x exists, there are integers t, t' such that $x = 6t + 1 = 15t' + 2$. Taking everything modulo $3 = \gcd(6, 15)$ we get $1 = 2 \pmod{3}$ which is false.

- (b) [10 points] Given four positive integers a_1, a_2, b_1 , and b_2 , all of which are no more than 2^n , give an algorithm running in polynomial time in n that returns `True` if there exists some x satisfying:

$$\begin{aligned}x &\equiv a_1 \pmod{b_1} \\x &\equiv a_2 \pmod{b_2}\end{aligned}$$

and `False` if no such x exists. Note that b_1 and b_2 need not necessarily be relatively prime.

Solution: Find the GCD of b_1 and b_2 using Euclid's algorithm. Let $g = \gcd(b_1, b_2)$. If $a_1 \equiv a_2 \pmod{g}$, return "True," otherwise return "False."

Problem 9. Sensible Subsets [30 points]

A set of integers S is called *k-sensible* if the sum of any k numbers in S is greater than or equal to every number in S .

- (a) [10 points] Given an array A of n distinct positive integers, prove that the maximum size *k-sensible* subset S of A contains consecutively valued numbers of A . In other words, if S contains a and b and $a < b$, then S must also contain every number x in A such that $a < x < b$. Assume that the maximum size *k-sensible* subset S of A is unique.

Solution: Denote the sum of the k smallest elements of S as t and the maximum of S to be m . First, we note that we can simplify the *k-sensible* condition to proving that $t \geq m$. This is because transitively, if $t \geq m$, t must also be greater than or equal to every number in S . If t is greater than or equal to every number in S , every sum of k numbers must be, since that sum will always be greater than t . Next, we note that if some number u between the minimum of S and m is in A but not S , we can simply add u to S to get a larger *k-sensible* subset since $s > m > u$ so S cannot be the largest. Thus, S must contain consecutively valued numbers of A .

- (b) [20 points] Give an algorithm that takes an array A and integer k and returns the maximum size k -sensible subset of A . Analyze your runtime in terms of n (you may assume $k \leq n$). Full credit will be given to algorithms that run in $O(n \log n)$.

Solution: Sort the numbers in increasing order. Compute the sum of every k length subarray in the sorted array by doing a "rolling sum". (Sum up the first k numbers, and then for each subsequent number, subtract the first and add on the new number. For each k length subarray sum, binary search for the greatest number in the array less than the sum. Return the longest length subarray found from the beginning of the subsequence sum to the number found by binary search.

Runtime is $O(n \log n)$ for sorting, $O(n)$ for rolling sum, and $O(n \log n)$ for n binary searches for overall $O(n \log n)$.

This problem can also alternatively be done without binary search by keeping track of the index of the largest number less than or equal to the subarray sums. Since this index only ever increases, it only iterates through the list once, giving us $O(n)$ time for everything but the sorting. A previous version of this problem made you come up with this $O(n)$ trick and use counting sort (the numbers were given as integers) for overall $O(n)$ runtime but we decided to be nice :).

Note that if the sum of the k smallest numbers beginning at an index is computed for every index without the rolling sum, this takes $O(k)$ time for each of the $O(n)$ problems for overall $O(n \log n + nk)$ runtime.

SCRATCH PAPER

SCRATCH PAPER