

1 Algorithmic Thinking

For each function $f(n)$ along the left side of the table, and for each function $g(n)$ across the top, write O , Ω , or Θ in the appropriate space, depending on whether $f(n) = O(g(n))$, $f(n) = \Omega(g(n))$, or $f(n) = \Theta(g(n))$. If more than one such relation holds between $f(n)$ and $g(n)$, write *only the strongest one* (which is the only answer considered correct).

The first row is a demo solution for $f(n) = n \log n$. The function \log denotes logarithm to the base two unless otherwise noted. $\binom{n}{2}$ denotes the “ n choose 2” symbol.

		$g(n)$		
		$n^{0.01}$	$n^{\log \log n}$	$n \log_{15} n$
$f(n)$	$n \log n$	Ω	O	Θ
	$(\log n)^{\log n}$			
	$2^{\sqrt{\log n}}$			
	$n \log \binom{n}{2}$			

2 Peak Finding

Design an algorithm that finds the maximum of a *rotated-sorted* array of n **distinct** elements in $O(\log n)$ time. A rotated-sorted array is formed by circularly shifting an array of integers arranged in increasing order (for example, $[9, 11, 2, 3, 6]$ is one such array.)

Note: You do not need to prove that your algorithm is correct, but you need to provide a **brief** analysis of its running time.

3 Document distance

Name an operation that was one of the biggest asymptotic “time-sinks” in the original (un-optimized) version of the document distance program. (There may be more than one correct answer; any one will do.) What running time was incurred by this “time sink”?

4 Insertion Sort and Merge Sort

True / False Running merge sort on an array of size n which is already correctly sorted takes $O(n)$ time.

True / False Insertion Sort makes more comparisons than Heap Sort on all inputs.

5 Master theorem

(Fall 14)

$T(n) = T(n/4) + T(3n/4) + \Theta(n)$ **(Spring 14)** What is the asymptotic runtime of an algorithm with the following recurrence:

$$T(n) = T\left(\frac{n}{3}\right) + T\left(\frac{n}{6}\right) + \Theta(n)$$

- (a) $\Theta(\log n)$
- (b) $\Theta(n)$
- (c) $\Theta(n \log n)$
- (d) $\Theta(n^2)$
- (e) $\Theta(n^2 \log n)$

6 Heaps and HeapSort

Suppose a binary max-heap H contains 80 distinct keys. How many distinct positions might contain the *smallest* element in H ?

- (a) 32
- (b) 40
- (c) 41
- (d) 42
- (e) 64

Which of the two algorithms {HeapSort, MergeSort}, implemented as described in class, is a better choice if **space** (memory usage) is the primary concern, rather than running time?

7 Binary Search Trees, BST Sort

(Spring 15)

Given a **sorted list** of n numbers (not necessarily integers), give an algorithm to construct a **balanced binary search tree** containing the same numbers in time $O(n)$, or argue why it's not possible.

Problem: Describe (in clear English or pseudocode) an $O(n)$ algorithm for balancing an arbitrary binary search tree, that is, for producing a new binary search tree with the same elements as the original one but with height $O(\log(n))$. Include a time complexity analysis.

8 AVL Trees, AVL Sort

(Fall 15) **Superconductivity Data Collection** (4 parts)

In a series of experiments on superconductivity, you are collecting resistance measurements that were observed at different temperatures. Each data point contains a temperature reading and the resistance measured at that temperature. For example, the data point $(10, 3)$ was observed at temperature 10 (kelvins) and has resistance 3 (nano-ohms). At any time during your experiments, you would like to be able to insert a new data point, and to find the data point with the minimum resistance that was observed at or below temperature t kelvins.

Create a data structure that will allow you to insert data points, $\text{INSERT}((t, r))$ where t is the temperature and r is the measured resistance in $O(\log n)$ runtime. Your data structure must also allow finding the data point with the minimum resistance at or below a given temperature t , $\text{FIND-MINIMUM-RESISTANCE}(t)$, in $O(\log n)$ runtime. You may assume that the set of data points is initially empty. The following table shows an example of a sequence of operations, along with their desired behaviors. **Clarification:** The temperatures of all data points are different.

operation	desired behavior
$\text{INSERT}((10, 3))$	update the set of data points to $\{(10, 3)\}$
$\text{INSERT}((2, 5))$	update the set of data points to $\{(10, 3), (2, 5)\}$
$\text{FIND-MINIMUM-RESISTANCE}(10)$	return the data point $(10, 3)$
$\text{FIND-MINIMUM-RESISTANCE}(5)$	return the data point $(2, 5)$
$\text{FIND-MINIMUM-RESISTANCE}(1)$	return None
$\text{INSERT}((4.5, 0.2))$	update the set of data points to $\{(10, 3), (2, 5), (4.5, 0.2)\}$
$\text{FIND-MINIMUM-RESISTANCE}(5)$	return the data point $(4.5, 0.2)$

- (a) What data structure that we learned in class could be useful here? Describe how this data structure can be used to store the data.
- (b) How can this data structure be augmented to enable efficient $\text{FIND-MINIMUM-RESISTANCE}(t)$ queries?
- (c) Describe how to insert into the data structure, $\text{INSERT}((t, r))$, while maintaining this augmentation in $O(\log n)$ time.
- (d) Describe how to perform $\text{FIND-MINIMUM-RESISTANCE}(t)$ in $O(\log n)$ time on the data structure you described above.

9 Sorting Lower Bounds

Problem Consider the following multi-dictionary problem. Let $A[1..n]$ be a fixed sorted array of distinct integers. Given an array $X[1..k]$, we want to find the position (if any) of each integer $X[i]$ in the array A . In other words, we want to compute an array $I[1..k]$ where for each i , either $I[i] = 0$ (so zero means none) or $A[I[i]] = X[i]$. Determine the exact asymptotic complexity of this problem, as a function of n and k in the binary decision tree model.

(Fall 10) We can sort 7 numbers with 10 comparisons.

10 Counting Sort, Radix Sort

Ben Bitdiddle modifies RADIX-SORT to use INSERTION-SORT to sort by digits, instead of COUNTING-SORT . Would the resulting algorithm still work correctly? What is the complexity of this new algorithm? The complexity of conventional RADIX-SORT on n numbers in base b with at most d digits is $O((n + b) \cdot d)$.

11 Hashing with Chaining

Given an array A of n integers and an integer k , detect if there is an entry $A[i]$ that is equal to one of the k previous entries $A[i - 1] \dots A[i - k]$. Your algorithm should run in time $O(n)$. You can assume you have access to a hash function which satisfies the simple uniform hashing assumption (SUHA).

Example: Given an array $A = [1, 3, 5, 7, 6, 5, 2]$ and $k = 4$, the algorithm should output YES since $A[3] = A[6] = 5$.

Note: You do not need to prove that your algorithm is correct, but you need to provide a **brief** analysis of its running time.

12 Table Doubling, Karp-Rabin

(Spring 16) Consider a modification to the table doubling procedure in which whenever the number of keys n is at least $\frac{m}{2}$, where m is the current size of the table, we set its new size m' to be $m' = 2m + \lfloor \sqrt{m} \rfloor$. If we never delete any elements from our table, then the resulting amortized overhead of this modified table doubling is still only $O(1)$ per each hash table operation.

(Fall 12) When using “Table Doubling” to maintain the size of a hash table (with insertions but no deletions), the size of the table grows exponentially with the number n of keys inserted.

(Fall 12) Recall the Rabin-Karp string-matching algorithm, which uses a rolling hash to search for a pattern of length m in a text of length n . Suppose it is run until the first match, if any, is found. Then the expected running time is $\Theta(nm)$.